

The argument for Decoupled Triton

Quinn Pham

Fused Multiply-Add Example

Fused Multiply-Add Example

Decoupled Triton

Fused Multiply-Add Example

Decoupled Triton

```
In a;  
In b;
```

Fused Multiply-Add Example

Decoupled Triton

```
In a;  
In b;  
SIn s;
```

Fused Multiply-Add Example

Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;
```

Fused Multiply-Add Example

Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;
```

Fused Multiply-Add Example

Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];
```

Fused Multiply-Add Example

Decoupled Triton

```
In a;
In b;
SIn s;
Func f;
Var x;
Var y;

f[x, y] = a[x, y] * s + b[x, y];

f.compile_to_kernel();
```

Fused Multiply-Add Example

Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.compile_to_kernel();
```

Triton Python

Fused Multiply-Add Example

Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.compile_to_kernel();
```

Triton Python

```
import triton  
import triton.language as tl
```

Fused Multiply-Add Example

Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.compile_to_kernel();
```

Triton Python

```
import triton  
import triton.language as tl  
  
@triton.jit  
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y):
```

Fused Multiply-Add Example

Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.compile_to_kernel();
```

Triton Python

```
import triton  
import triton.language as tl  
  
@triton.jit  
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y):  
    for x_iter in range(0, x, 1):  
        for y_iter in range(0, y, 1):
```

Fused Multiply-Add Example

Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.compile_to_kernel();
```

Triton Python

```
import triton  
import triton.language as tl  
  
@triton.jit  
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y):  
    for x_iter in range(0, x, 1):  
        for y_iter in range(0, y, 1):  
            a = tl.load(a_ptr + x_iter * y + y_iter)  
            b = tl.load(b_ptr + x_iter * y + y_iter)
```

Fused Multiply-Add Example

Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.compile_to_kernel();
```

Triton Python

```
import triton  
import triton.language as tl  
  
@triton.jit  
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y):  
    for x_iter in range(0, x, 1):  
        for y_iter in range(0, y, 1):  
            a = tl.load(a_ptr + x_iter * y + y_iter)  
            b = tl.load(b_ptr + x_iter * y + y_iter)  
            f = a * s + b
```

Fused Multiply-Add Example

Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.compile_to_kernel();
```

Triton Python

```
import triton  
import triton.language as tl  
  
@triton.jit  
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y):  
    for x_iter in range(0, x, 1):  
        for y_iter in range(0, y, 1):  
            a = tl.load(a_ptr + x_iter * y + y_iter)  
            b = tl.load(b_ptr + x_iter * y + y_iter)  
            f = a * s + b  
            tl.store(f_ptr + x_iter * y + y_iter, f)
```

Fused Multiply-Add Example

Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.compile_to_kernel();
```

Triton Python

```
import triton  
import triton.language as tl  
  
@triton.jit  
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y):  
    for x_iter in range(0, x, 1):  
        for y_iter in range(0, y, 1):  
            a = tl.load(a_ptr + x_iter * y + y_iter)  
            b = tl.load(b_ptr + x_iter * y + y_iter)  
            f = a * s + b  
            tl.store(f_ptr + x_iter * y + y_iter, f)  
  
...  
f_grid = (1,)  
f_kernel[f_grid](a, b, s, f, x, y)
```


Programmability

What advantages does defining kernels in Decoupled Triton have over writing Triton Kernels by hand?

Manually modifying an existing Triton kernel according to a schedule is non-trivial due to index remappings, offset calculations, and other low-level details


Blocking

Blocking





Blocking

```
f.block(x:2, y:4)
```





Blocking

```
f.block(x:2, y:4)
```





Blocking

```
f.block(x:2, y:4)
```





Blocking

```
f.block(x:2, y:4)
```





Blocking

```
f.block(x:2, y:4)
```





Blocking

`f.block(x:2, y:4)`




Blocking

`f.block(x:2, y:4)`



Blocking


`f.block(x:2, y:4)`



0, 0	0, 1	0, 2	0, 3
1, 0	1, 1	1, 2	1, 3
2, 0	2, 1	2, 2	2, 3
3, 0	3, 1	3, 2	3, 3
4, 0	4, 1	4, 2	4, 3
5, 0	5, 1	5, 2	5, 3
6, 0	6, 1	6, 2	6, 3
7, 0	7, 1	7, 2	7, 3

Blocking

`f.block(x:2, y:4)`



0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
16	17	18	19
20	21	22	23
24	25	26	27
28	29	30	31

Blocking in Triton

Blocking in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)
    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for x_iter in range(0, x_BLOCK_SIZE, 1):
        for y_iter in range(0, y_BLOCK_SIZE, 1):
            a = tl.load(a_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            b = tl.load(b_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            f = a * s + b
            tl.store(f_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter, f)

...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4)
```

Blocking in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)
    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for x_iter in range(0, x_BLOCK_SIZE, 1):
        for y_iter in range(0, y_BLOCK_SIZE, 1):
            a = tl.load(a_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            b = tl.load(b_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            f = a * s + b
            tl.store(f_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter, f)
...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4)
```

Blocking in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)
    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for x_iter in range(0, x_BLOCK_SIZE, 1):
        for y_iter in range(0, y_BLOCK_SIZE, 1):
            a = tl.load(a_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            b = tl.load(b_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            f = a * s + b
            tl.store(f_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter, f)

...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4)
```

Blocking in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)
    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for x_iter in range(0, x_BLOCK_SIZE, 1):
        for y_iter in range(0, y_BLOCK_SIZE, 1):
            a = tl.load(a_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            b = tl.load(b_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            f = a * s + b
            tl.store(f_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter, f)

...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4)
```

Blocking in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)
    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for x_iter in range(0, x_BLOCK_SIZE, 1):
        for y_iter in range(0, y_BLOCK_SIZE, 1):
            a = tl.load(a_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            b = tl.load(b_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            f = a * s + b
            tl.store(f_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter, f)
    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4)
```

Blocking in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)
    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for x_iter in range(0, x_BLOCK_SIZE, 1):
        for y_iter in range(0, y_BLOCK_SIZE, 1):
            a = tl.load(a_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            b = tl.load(b_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            f = a * s + b
            tl.store(f_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter, f)

...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4)
```

Blocking in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)
    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for x_iter in range(0, x_BLOCK_SIZE, 1):
        for y_iter in range(0, y_BLOCK_SIZE, 1):
            a = tl.load(a_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            b = tl.load(b_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            f = a * s + b
            tl.store(f_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter, f)

...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4)
```

Blocking in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)
    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for x_iter in range(0, x_BLOCK_SIZE, 1):
        for y_iter in range(0, y_BLOCK_SIZE, 1):
            a = tl.load(a_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            b = tl.load(b_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter)
            f = a * s + b
            tl.store(f_ptr + (x_block_start + x_iter) * y + y_block_start + y_iter, f)
    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4)
```

Blocking in Decoupled Triton

Blocking in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.compile_to_kernel();
```

Blocking in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];
```

```
f.block(x:2, y:4);  
f.compile_to_kernel();
```

Blocking in Decoupled Triton


```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.compile_to_kernel();
```

Tensorization

Tensorization

```
f.block(x:2, y:4)
```


```
f.tensorize(x:2, y:2)
```



Tensorization

```
f.block(x:2, y:4)
```


```
f.tensorize(x:2, y:2)
```



Tensorization

```
f.block(x:2, y:4)
```


```
f.tensorize(x:2, y:2)
```



Tensorization

```
f.block(x:2, y:4)
```


```
f.tensorize(x:2, y:2)
```



Tensorization

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```



Tensorization in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)
    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[ :, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[ :, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[ :, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2)
```

Tensorization in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)
    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[ :, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[ :, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[ :, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2)
```

Tensorization in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)
    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2)
```

Tensorization in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)
    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[ :, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[ :, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[ :, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2)
```

Tensorization in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)
    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[ :, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[ :, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[ :, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2)
```

Tensorization in Decoupled Triton

Tensorization in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.tensorize(x: 2, y:2);  
f.compile_to_kernel();
```

Tensorization in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.tensorize(x: 2, y:2);  
f.compile_to_kernel();
```

Tensorization in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.tensorize(x: 2, y:2);  
f.compile_to_kernel();
```


Grouping

Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```




Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```




Grouping

`f.block(x:2, y:4)`

`f.tensorize(x:2, y:2)`

`f.group(x:4, y:2)`




Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```




0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
16	17	18	19
20	21	22	23
24	25	26	27
28	29	30	31

Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```





Grouping

`f.block(x:2, y:4)`

`f.tensorize(x:2, y:2)`

`f.group(x:4, y:2)`




Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```




Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```




Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```



Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```




0	1		
2	3		
4	5		
6	7		

Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```



0	1	8	9
2	3	10	11
4	5	12	13
6	7	14	15
16	17	24	25
18	19	26	27
20	21	28	29
22	23	30	31


Why group?

Why group?

An operation for which a **group** of the result has better data locality than the default row-major ordering?



Why group?

An operation for which a **group** of the result has better data locality than the default row-major ordering?




Why group?

An operation for which a **group** of the result has better data locality than the default row-major ordering?



Matrix Multiplication




Matrix Multiplication


\times

$=$



Matrix Multiplication




\times




$=$



\times



$=$



Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :])
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], a.dtype)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], b.dtype)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ... # Group Remapping

    pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
    num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
    group_id = pid // num_blocks_in_group
    group_offset = pid % num_blocks_in_group
    y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
    x_group_id = group_id // y_num_group_ids
    y_group_id = group_id % y_num_group_ids
    x_group_offset = group_offset // y_GROUP_SIZE
    y_group_offset = group_offset % y_GROUP_SIZE
    x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
    y_pid = y_group_id * y_GROUP_SIZE + y_group_offset

f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE'])),
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Group Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
16	17	18	19
20	21	22	23
24	25	26	27
28	29	30	31

Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE'])),
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Group Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
16	17	18	19
20	21	22	23
24	25	26	27
28	29	30	31

Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Group Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```



Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Group Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```



8

Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)


    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :, f])
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :, f], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Group Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```



Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

...

Group Remapping

```
pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```

The diagram shows a 2x2 grid of four white squares. The bottom-left square is highlighted with a green border and contains the number '2'. This visualizes the grouping of threads into 2x2 blocks, where each block is assigned a unique identifier (2 in this case).

Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ... # Group Remapping

    pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
    num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
    group_id = pid // num_blocks_in_group
    group_offset = pid % num_blocks_in_group
    y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
    x_group_id = group_id // y_num_group_ids
    y_group_id = group_id % y_num_group_ids
    x_group_offset = group_offset // y_GROUP_SIZE
    y_group_offset = group_offset % y_GROUP_SIZE
    x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
    y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
    f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
    f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Group Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```



Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)


    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Group Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```



Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], a.dtype)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], b.dtype)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

... # Group Remapping

```
pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```



Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Group Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```



2

Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)


    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], a.dtype)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :, b.dtype])
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :, b.dtype], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Group Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```



Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)


    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], a.dtype)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], b.dtype)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

... # Group Remapping

```
pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```



1, 0

Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ... # Group Remapping

    pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
    num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
    group_id = pid // num_blocks_in_group
    group_offset = pid % num_blocks_in_group
    y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
    x_group_id = group_id // y_num_group_ids
    y_group_id = group_id % y_num_group_ids
    x_group_offset = group_offset // y_GROUP_SIZE
    y_group_offset = group_offset % y_GROUP_SIZE
    x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
    y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
    f = a * s + b
    tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
             y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Group Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```



Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Group Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```

	1,1

Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)


    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], a.dtype)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], b.dtype)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

... # Group Remapping

```
pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```



Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)


    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

... # Group Remapping

```
pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
group_id = pid // num_blocks_in_group
group_offset = pid % num_blocks_in_group
y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
x_group_id = group_id // y_num_group_ids
y_group_id = group_id % y_num_group_ids
x_group_offset = group_offset // y_GROUP_SIZE
y_group_offset = group_offset % y_GROUP_SIZE
x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
```



Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ... # Group Remapping

    pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
    num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
    group_id = pid // num_blocks_in_group
    group_offset = pid % num_blocks_in_group
    y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
    x_group_id = group_id // y_num_group_ids
    y_group_id = group_id % y_num_group_ids
    x_group_offset = group_offset // y_GROUP_SIZE
    y_group_offset = group_offset % y_GROUP_SIZE
    x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
    y_pid = y_group_id * y_GROUP_SIZE + y_group_offset
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Group Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], a.dtype)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], b.dtype)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ... # Group Remapping

    pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
    num_blocks_in_group = x_GROUP_SIZE * y_GROUP_SIZE
    group_id = pid // num_blocks_in_group
    group_offset = pid % num_blocks_in_group
    y_num_group_ids = y // y_BLOCK_SIZE // y_GROUP_SIZE
    x_group_id = group_id // y_num_group_ids
    y_group_id = group_id % y_num_group_ids
    x_group_offset = group_offset // y_GROUP_SIZE
    y_group_offset = group_offset % y_GROUP_SIZE
    x_pid = x_group_id * x_GROUP_SIZE + x_group_offset
    y_pid = y_group_id * y_GROUP_SIZE + y_group_offset

```

Grouping in Decoupled Triton

Grouping in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.tensorize(x: 2, y:2);  
f.group(x:4, y:2);  
f.compile_to_kernel();
```

Grouping in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.tensorize(x: 2, y:2);  
f.group(x:4, y:2);  
f.compile_to_kernel();
```

Grouping in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.tensorize(x: 2, y:2);  
f.group(x:4, y:2);  
f.compile_to_kernel();
```


Dilation

Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```




Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```





Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```




Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```





0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
16	17	18	19
20	21	22	23
24	25	26	27
28	29	30	31

Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```





Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```





Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```





Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```





Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```





Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```





Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```





Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```




Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```




0	4	1	5
2	6	3	7

Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```




0	4	1	5
8			
2	6	3	7

Dilation

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.dilate(x:4, y:2)
```



0	4	1	5
8	12	9	13
16	20	17	21
24	28	25	29
2	6	3	7
10	14	11	15
18	22	19	23
26	30	27	31


Why dilate?

Why dilate?

An operation for which a **dilation** of the result has better data locality than the default row-major ordering?



Why dilate?

An operation for which a **dilation** of the result has better data locality than the default row-major ordering?



Why dilate?

An operation for which a **dilation** of the result has better data locality than the default row-major ordering?



Why dilate?

An operation for which a **dilation** of the result has better data locality than the default row-major ordering?

Does one exist?

Dilation in Triton

Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::])
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::])
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::])
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::])
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::])
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::])
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
    num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                      (y // y_BLOCK_SIZE) // y_DILATE_SIZE
    dilate_id = pid % num_dilate_ids
    dilate_offset = pid // num_dilate_ids
    y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
    x_dilate_id = dilate_id // y_num_dilate_ids
    y_dilate_id = dilate_id % y_num_dilate_ids
    x_dilate_offset = dilate_offset // y_DILATE_SIZE
    y_dilate_offset = dilate_offset % y_DILATE_SIZE
    x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
    y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset

    a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
               y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
    b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
               y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
    f = a * s + b
    tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
             y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Dilation Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE * (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
16	17	18	19
20	21	22	23
24	25	26	27
28	29	30	31

Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

... # Dilation Remapping

```
pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                  (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
16	17	18	19
20	21	22	23
24	25	26	27
28	29	30	31

Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping


    pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
    num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                      (y // y_BLOCK_SIZE) // y_DILATE_SIZE
    dilate_id = pid % num_dilate_ids
    dilate_offset = pid // num_dilate_ids
    y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
    x_dilate_id = dilate_id // y_num_dilate_ids
    y_dilate_id = dilate_id % y_num_dilate_ids
    x_dilate_offset = dilate_offset // y_DILATE_SIZE
    y_dilate_offset = dilate_offset % y_DILATE_SIZE
    x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
    y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset

    f = a * s + b
    tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
             y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Dilation Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                  (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```




Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
    num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                      (y // y_BLOCK_SIZE) // y_DILATE_SIZE
    dilate_id = pid % num_dilate_ids
    dilate_offset = pid // num_dilate_ids
    y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
    x_dilate_id = dilate_id // y_num_dilate_ids
    y_dilate_id = dilate_id % y_num_dilate_ids
    x_dilate_offset = dilate_offset // y_DILATE_SIZE
    y_dilate_offset = dilate_offset % y_DILATE_SIZE
    x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
    y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
    f = a * s + b
    tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
             y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```



4

Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)


    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Dilation Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE * (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```



Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)


    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Dilation Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE * (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```



Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)


    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Dilation Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                  (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```



Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Dilation Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                  (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```

0	

Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Dilation Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                  (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```



Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Dilation Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                  (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```



Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)


    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Dilation Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                  (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```



Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)


    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Dilation Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                  (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```



Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Dilation Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                  (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```



Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Dilation Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                  (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```

0,0	

Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)


    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

```
... # Dilation Remapping

pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                  (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```



Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)


    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

... # Dilation Remapping

```
pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                  (y // y_BLOCK_SIZE) // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset
```



Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
    num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                      (y // y_BLOCK_SIZE) // y_DILATE_SIZE
    dilate_id = pid % num_dilate_ids
    dilate_offset = pid // num_dilate_ids
    y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
    x_dilate_id = dilate_id // y_num_dilate_ids
    y_dilate_id = dilate_id % y_num_dilate_ids
    x_dilate_offset = dilate_offset // y_DILATE_SIZE
    y_dilate_offset = dilate_offset % y_DILATE_SIZE
    x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
    y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset

    a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
               y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
    b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
               y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
    f = a * s + b
    tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
             y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

Dilation in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], a)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], b)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[::, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, ::], f)

    ... # Dilation Remapping

    pid = x_pid * (y // y_BLOCK_SIZE) + y_pid
    num_dilate_ids = (x // x_BLOCK_SIZE) // x_DILATE_SIZE *
                     (y // y_BLOCK_SIZE) // y_DILATE_SIZE
    dilate_id = pid % num_dilate_ids
    dilate_offset = pid // num_dilate_ids
    y_num_dilate_ids = (y // y_BLOCK_SIZE) // y_DILATE_SIZE
    x_dilate_id = dilate_id // y_num_dilate_ids
    y_dilate_id = dilate_id % y_num_dilate_ids
    x_dilate_offset = dilate_offset // y_DILATE_SIZE
    y_dilate_offset = dilate_offset % y_DILATE_SIZE
    x_pid = x_dilate_id * x_DILATE_SIZE + x_dilate_offset
    y_pid = y_dilate_id * y_DILATE_SIZE + y_dilate_offset

f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2)
```

Dilation in Decoupled Triton

Dilation in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.tensorize(x: 2, y:2);  
f.dilate(x:4, y:2);  
f.compile_to_kernel();
```

Dilation in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.tensorize(x: 2, y:2);  
f.dilate(x:4, y:2);  
f.compile_to_kernel();
```

Dilation in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.tensorize(x: 2, y:2);  
f.dilate(x:4, y:2);  
f.compile_to_kernel();
```


Dilation with Grouping

f.block(x:2, y:4)

f.tensorize(x:2, y:2)

f.group(x:4, y:2)

f.dilate(x:2, y:1)





Dilation with Grouping

`f.block(x:2, y:4)`

`f.tensorize(x:2, y:2)`

`f.group(x:4, y:2)`

`f.dilate(x:2, y:1)`





Dilation with Grouping

`f.block(x:2, y:4)`

`f.tensorize(x:2, y:2)`

`f.group(x:4, y:2)`

`f.dilate(x:2, y:1)`





Dilation with Grouping

`f.block(x:2, y:4)`

`f.tensorize(x:2, y:2)`

`f.group(x:4, y:2)`

`f.dilate(x:2, y:1)`





Dilation with Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```

```
f.dilate(x:2, y:1)
```




Dilation with Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```

```
f.dilate(x:2, y:1)
```




Dilation with Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```

```
f.dilate(x:2, y:1)
```




Dilation with Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```

```
f.dilate(x:2, y:1)
```




Dilation with Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```

```
f.dilate(x:2, y:1)
```



0	1		
4			
2	3		


Dilation with Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```

```
f.dilate(x:2, y:1)
```



0	1		
4	5		
2	3		
6	7		


Dilation with Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```

```
f.dilate(x:2, y:1)
```



0	1	8	9
4	5	12	13
2	3	10	11
6	7	14	15
16	17	24	25
20	21	28	29
18	19	26	27
22	23	30	31

Dilation with Grouping in Triton

Dilation with Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr, x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Grouping Remapping
    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2, 2, 1)
```

```
... # Dilation Remapping

    x_pid = x_pid % x_GROUP_SIZE
    y_pid = y_pid % y_GROUP_SIZE
    pid = x_pid * y_GROUP_SIZE + y_pid
    num_dilate_ids = x_GROUP_SIZE // x_DILATE_SIZE * y_GROUP_SIZE // y_DILATE_SIZE
    dilate_id = pid % num_dilate_ids
    dilate_offset = pid // num_dilate_ids
    y_num_dilate_ids = y_GROUP_SIZE // y_DILATE_SIZE
    x_dilate_id = dilate_id // y_num_dilate_ids
    y_dilate_id = dilate_id % y_num_dilate_ids
    x_dilate_offset = dilate_offset // y_DILATE_SIZE
    y_dilate_offset = dilate_offset % y_DILATE_SIZE
    x_pid = x_group_id * x_GROUP_SIZE + (x_dilate_id * x_DILATE_SIZE + x_dilate_offset)
    y_pid = y_group_id * y_GROUP_SIZE + (y_dilate_id * y_DILATE_SIZE + y_dilate_offset)
```

Dilation with Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr, x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Grouping Remapping
    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2, 2, 1)
```

```
... # Dilation Remapping

x_pid = x_pid % x_GROUP_SIZE
y_pid = y_pid % y_GROUP_SIZE
pid = x_pid * y_GROUP_SIZE + y_pid
num_dilate_ids = x_GROUP_SIZE // x_DILATE_SIZE * y_GROUP_SIZE // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = y_GROUP_SIZE // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_group_id * x_GROUP_SIZE + (x_dilate_id * x_DILATE_SIZE + x_dilate_offset)
y_pid = y_group_id * y_GROUP_SIZE + (y_dilate_id * y_DILATE_SIZE + y_dilate_offset)
```

Dilation with Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr, x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Grouping Remapping
    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2, 2, 1)
```

... # Dilation Remapping

```
x_pid = x_pid % x_GROUP_SIZE
y_pid = y_pid % y_GROUP_SIZE
pid = x_pid * y_GROUP_SIZE + y_pid
num_dilate_ids = x_GROUP_SIZE // x_DILATE_SIZE * y_GROUP_SIZE // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = y_GROUP_SIZE // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_group_id * x_GROUP_SIZE + (x_dilate_id * x_DILATE_SIZE + x_dilate_offset)
y_pid = y_group_id * y_GROUP_SIZE + (y_dilate_id * y_DILATE_SIZE + y_dilate_offset)
```

Dilation with Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr, x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Grouping Remapping
    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2, 2, 1)
```

```
... # Dilation Remapping

x_pid = x_pid % x_GROUP_SIZE
y_pid = y_pid % y_GROUP_SIZE
pid = x_pid * y_GROUP_SIZE + y_pid
num_dilate_ids = x_GROUP_SIZE // x_DILATE_SIZE * y_GROUP_SIZE // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = y_GROUP_SIZE // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_group_id * x_GROUP_SIZE + (x_dilate_id * x_DILATE_SIZE + x_dilate_offset)
y_pid = y_group_id * y_GROUP_SIZE + (y_dilate_id * y_DILATE_SIZE + y_dilate_offset)
```

Dilation with Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr, x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Grouping Remapping
    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2, 2, 1)
```

```
... # Dilation Remapping

x_pid = x_pid % x_GROUP_SIZE
y_pid = y_pid % y_GROUP_SIZE
pid = x_pid * y_GROUP_SIZE + y_pid
num_dilate_ids = x_GROUP_SIZE // x_DILATE_SIZE * y_GROUP_SIZE // y_DILATE_SIZE
dilate_id = pid % num_dilate_ids
dilate_offset = pid // num_dilate_ids
y_num_dilate_ids = y_GROUP_SIZE // y_DILATE_SIZE
x_dilate_id = dilate_id // y_num_dilate_ids
y_dilate_id = dilate_id % y_num_dilate_ids
x_dilate_offset = dilate_offset // y_DILATE_SIZE
y_dilate_offset = dilate_offset % y_DILATE_SIZE
x_pid = x_group_id * x_GROUP_SIZE + (x_dilate_id * x_DILATE_SIZE + x_dilate_offset)
y_pid = y_group_id * y_GROUP_SIZE + (y_dilate_id * y_DILATE_SIZE + y_dilate_offset)
```

Dilation with Grouping in Triton

```
...
@triton.jit
def f_kernel(a_ptr, b_ptr, s, f_ptr, x, y, x_BLOCK_SIZE: tl.constexpr, y_BLOCK_SIZE: tl.constexpr, y_TENSOR_SIZE: tl.constexpr
             x_GROUP_SIZE: tl.constexpr, y_GROUP_SIZE: tl.constexpr, x_DILATE_SIZE: tl.constexpr, y_DILATE_SIZE: tl.constexpr):
    x_pid = tl.program_id(0)
    y_pid = tl.program_id(1)

    ... # Grouping Remapping
    ... # Dilation Remapping

    x_block_start = x_pid * x_BLOCK_SIZE
    y_block_start = y_pid * y_BLOCK_SIZE
    for y_iter in range(0, y_BLOCK_SIZE, y_TENSOR_SIZE):
        a = tl.load(a_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        b = tl.load(b_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                   y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], tl.float32)
        f = a * s + b
        tl.store(f_ptr + (x_block_start + tl.arange(0, x_BLOCK_SIZE)[:, None]) * y +
                 y_block_start + y_iter + tl.arange(0, y_TENSOR_SIZE)[None, :], f)

    ...
f_grid = lambda meta: (triton.cdiv(x, meta['x_BLOCK_SIZE']), triton.cdiv(y, meta['y_BLOCK_SIZE']),)
f_kernel[f_grid](a, b, s, f, x, y, 2, 4, 2, 4, 2, 2, 1)
```

```
... # Dilation Remapping

    x_pid = x_pid % x_GROUP_SIZE
    y_pid = y_pid % y_GROUP_SIZE
    pid = x_pid * y_GROUP_SIZE + y_pid
    num_dilate_ids = x_GROUP_SIZE // x_DILATE_SIZE * y_GROUP_SIZE // y_DILATE_SIZE
    dilate_id = pid % num_dilate_ids
    dilate_offset = pid // num_dilate_ids
    y_num_dilate_ids = y_GROUP_SIZE // y_DILATE_SIZE
    x_dilate_id = dilate_id // y_num_dilate_ids
    y_dilate_id = dilate_id % y_num_dilate_ids
    x_dilate_offset = dilate_offset // y_DILATE_SIZE
    y_dilate_offset = dilate_offset % y_DILATE_SIZE
    x_pid = x_group_id * x_GROUP_SIZE + (x_dilate_id * x_DILATE_SIZE + x_dilate_offset)
    y_pid = y_group_id * y_GROUP_SIZE + (y_dilate_id * y_DILATE_SIZE + y_dilate_offset)
```

Dilation in Decoupled Triton

Dilation in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.tensorize(x: 2, y:2);  
f.group(x:4, y:2);  
f.dilate(x:2, y:1);  
f.compile_to_kernel();
```

Dilation in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.tensorize(x: 2, y:2);  
f.group(x:4, y:2);  
f.dilate(x:2, y:1);  
f.compile_to_kernel();
```

Dilation in Decoupled Triton

```
In a;  
In b;  
SIn s;  
Func f;  
Var x;  
Var y;  
  
f[x, y] = a[x, y] * s + b[x, y];  
  
f.block(x:2, y:4);  
f.tensorize(x: 2, y:2);  
f.group(x:4, y:2);  
f.dilate(x:2, y:1);  
f.compile_to_kernel();
```


Aggregate and Sequentialize

Aggregate and Sequentialize

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.aggregate_and_sequentialize(4)
```



Aggregate and Sequentialize

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

f.aggregate_and_sequentialize(4)

Aggregate and Sequentialize

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.aggregate_and_sequentialize(4)
```


0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15
16	17	18	19
20	21	22	23
24	25	26	27
28	29	30	31

Aggregate and Sequentialize

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.aggregate_and_sequentialize(4)
```




Aggregate and Sequentialize

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.aggregate_and_sequentialize(4)
```




Aggregate and Sequentialize

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.aggregate_and_sequentialize(4)
```





Aggregate and Sequentialize

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.aggregate_and_sequentialize(4)
```





Aggregate and Sequentialize

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.aggregate_and_sequentialize(4)
```




Aggregate and Sequentialize

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.aggregate_and_sequentialize(4)
```




Aggregate and Sequentialize

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

f.aggregate_and_sequentialize(4)




Aggregate and Sequentialize

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

f.aggregate_and_sequentialize(4)




Aggregate and Sequentialize

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.aggregate_and_sequentialize(4)
```



0	0	0	0
1	1	1	1
2	2	2	2
3	3	3	3
4	4	4	4
5	5	5	5
6	6	6	6
7	7	7	7


Aggregate and Sequentialize with Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```

```
f.aggregate_and_sequentialize(4)
```



Aggregate and Sequentialize with Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```

```
f.aggregate_and_sequentialize(4)
```

0	1	8	9
2	3	10	11
4	5	12	13
6	7	14	15
16	17	24	25
18	19	26	27
20	21	28	29
22	23	30	31

Aggregate and Sequentialize with Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```

```
f.aggregate_and_sequentialize(4)
```



0	1	8	9
2	3	10	11
4	5	12	13
6	7	14	15
16	17	24	25
18	19	26	27
20	21	28	29
22	23	30	31


Aggregate and Sequentialize with Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```

```
f.aggregate_and_sequentialize(4)
```




Aggregate and Sequentialize with Grouping

```
f.block(x:2, y:4)
```

```
f.tensorize(x:2, y:2)
```

```
f.group(x:4, y:2)
```


```
f.aggregate_and_sequentialize(4)
```



0	0	2	2
0	0	2	2
1	1	3	3
1	1	3	3
4	4	6	6
4	4	6	6
5	5	7	7
5	5	7	7

Aggregate and Sequentialize with Dilation

```
f.block(x:2, y:4)
f.tensorize(x:2, y:2)
f.dilate(x:4, y:2)
f.aggregate_and_sequentialize(4)
```



Aggregate and Sequentialize with Dilation

```
f.block(x:2, y:4)
f.tensorize(x:2, y:2)
f.dilate(x:4, y:2)
f.aggregate_and_sequentialize(4)
```

0	4	1	5
8	12	9	13
16	20	17	21
24	28	25	29
2	6	3	7
10	14	11	15
18	22	19	23
26	30	27	31

Aggregate and Sequentialize with Dilation





```
f.block(x:2, y:4)  
f.tensorize(x:2, y:2)  
f.dilate(x:4, y:2)  
f.aggregate_and_sequentialize(4)
```



0	4	1	5
8	12	9	13
16	20	17	21
24	28	25	29
2	6	3	7
10	14	11	15
18	22	19	23
26	30	27	31


Aggregate and Sequentialize with Dilation

```
f.block(x:2, y:4)  
f.tensorize(x:2, y:2)  
f.dilate(x:4, y:2)  
f.aggregate_and_sequentialize(4)
```



Aggregate and Sequentialize with Dilation

```
f.block(x:2, y:4)  
f.tensorize(x:2, y:2)  
f.dilate(x:4, y:2)  
f.aggregate_and_sequentialize(4)
```



0	1	0	1
2	3	2	3
4	5	4	5
6	7	6	7
0	1	0	1
2	3	2	3
4	5	4	5
6	7	6	7

Fuse at

Fuse at

```
In a;  
In b;  
SIn s;  
Func f;  
Func g;  
Var x;  
Var y;
```

```
g[x, y] = s;  
f[x, y] = a[x, y] * g[x, y] + b[x, y];
```

```
f.fuse_at(g, x);  
f.compile_to_kernel();
```

Fuse the computation of a func into the computation of another func at a particular loop level.

Schedule Directive Recap

Schedule Directive Recap

`compile_to_kernel()`

Compiles a Triton kernel

Schedule Directive Recap

compile_to_kernel()	Compiles a Triton kernel
block()	Enables program blocking

Schedule Directive Recap

<code>compile_to_kernel()</code>	Compiles a Triton kernel
<code>block()</code>	Enables program blocking
<code>tensorize()</code>	Enables SIMD operations within a program

Schedule Directive Recap

<code>compile_to_kernel()</code>	Compiles a Triton kernel
<code>block()</code>	Enables program blocking
<code>tensorize()</code>	Enables SIMD operations within a program
<code>group()</code>	Remaps program IDs to compute the result in a grouped order

Schedule Directive Recap

<code>compile_to_kernel()</code>	Compiles a Triton kernel
<code>block()</code>	Enables program blocking
<code>tensorize()</code>	Enables SIMD operations within a program
<code>group()</code>	Remaps program IDs to compute the result in a grouped order
<code>dilate()</code>	Remaps program IDs to compute the result in a dilated order

Schedule Directive Recap

<code>compile_to_kernel()</code>	Compiles a Triton kernel
<code>block()</code>	Enables program blocking
<code>tensorize()</code>	Enables SIMD operations within a program
<code>group()</code>	Remaps program IDs to compute the result in a grouped order
<code>dilate()</code>	Remaps program IDs to compute the result in a dilated order
<code>aggregate_and_sequentialize()</code>	Sequentializes block computation of nearby programs

Schedule Directive Recap

<code>compile_to_kernel()</code>	Compiles a Triton kernel
<code>block()</code>	Enables program blocking
<code>tensorize()</code>	Enables SIMD operations within a program
<code>group()</code>	Remaps program IDs to compute the result in a grouped order
<code>dilate()</code>	Remaps program IDs to compute the result in a dilated order
<code>aggregate_and_sequentialize()</code>	Sequentializes block computation of nearby programs
<code>fuse_at()</code>	Fuses a function into another at a certain loop level

Performance

Performance

Can Decoupled Triton generate Triton kernels with performance that is **at least** equivalent to that of manually written Triton kernels?

Performance

Can Decoupled Triton generate Triton kernels with performance that is **at least** equivalent to that of manually written Triton kernels?

Yes, with the existing scheduling directives Decoupled Triton can generate competitive Triton kernels for many common ML ops.

Performance

Can Decoupled Triton generate Triton kernels with performance that is **at least** equivalent to that of manually written Triton kernels?

Yes, with the existing scheduling directives Decoupled Triton can generate competitive Triton kernels for many common ML ops.

[Liger-Kernel / src / liger_kernel / ops / rms_norm.py](#)

[Liger-Kernel / src / liger_kernel / ops / layer_norm.py](#)

Performance

Can Decoupled Triton generate Triton kernels with performance that is **at least** equivalent to that of manually written Triton kernels?

Yes, with the existing scheduling directives Decoupled Triton can generate competitive Triton kernels for many common ML ops.

[Liger-Kernel / src / liger_kernel / ops / rms_norm.py](#)

[Liger-Kernel / src / liger_kernel / ops / layer_norm.py](#)

[Liger-Kernel / src / liger_kernel / ops / jsd.py](#)

Performance

Can Decoupled Triton generate Triton kernels with performance that is **at least** equivalent to that of manually written Triton kernels?

Yes, with the existing scheduling directives Decoupled Triton can generate competitive Triton kernels for many common ML ops.

[Liger-Kernel / src / liger_kernel / ops / rms_norm.py](#)

[Liger-Kernel / src / liger_kernel / ops / layer_norm.py](#)

[Liger-Kernel / src / liger_kernel / ops / jsd.py](#)

[Liger-Kernel / src / liger_kernel / ops / swiglu.py](#)

Performance

Can Decoupled Triton generate Triton kernels with performance that is **at least** equivalent to that of manually written Triton kernels?

Yes, with the existing scheduling directives Decoupled Triton can generate competitive Triton kernels for many common ML ops.

[Liger-Kernel / src / liger_kernel / ops / rms_norm.py](#)

[Liger-Kernel / src / liger_kernel / ops / layer_norm.py](#)

[Liger-Kernel / src / liger_kernel / ops / jsd.py](#)

[Liger-Kernel / src / liger_kernel / ops / swiglu.py](#)

[Liger-Kernel / src / liger_kernel / ops / geglu.py](#)

Performance

Can Decoupled Triton generate Triton kernels with performance that is **at least** equivalent to that of manually written Triton kernels?

Yes, with the existing scheduling directives Decoupled Triton can generate competitive Triton kernels for many common ML ops.

[Liger-Kernel / src / liger_kernel / ops / layer_norm.py](#)[Liger-Kernel / src / liger_kernel / ops / jsd.py](#)[Liger-Kernel / src / liger_kernel / ops / swiglu.py](#)[Liger-Kernel / src / liger_kernel / ops / geglu.py](#)

Performance

Can Decoupled Triton generate Triton kernels with performance that is **at least** equivalent to that of manually written Triton kernels?

Yes, with the existing scheduling directives Decoupled Triton can generate competitive Triton kernels for many common ML ops.

[Liger-Kernel / src / liger_kernel / ops / rms_norm.py](#)

[triton / python / tutorials / 01-vector-add.py](#)

[Liger-Kernel / src / liger_kernel / ops / layer_norm.py](#)

[triton / python / tutorials / 03-matrix-multiplication.py](#)

[Liger-Kernel / src / liger_kernel / ops / jsd.py](#)

[Liger-Kernel / src / liger_kernel / ops / swiglu.py](#)

[Liger-Kernel / src / liger_kernel / ops / geglu.py](#)

Performance

Can Decoupled Triton generate Triton kernels with performance that is **at least** equivalent to that of manually written Triton kernels?

Yes, with the existing scheduling directives Decoupled Triton can generate competitive Triton kernels for many common ML ops.

[Liger-Kernel / src / liger_kernel / ops / rms_norm.py](#)

[Liger-Kernel / src / liger_kernel / ops / layer_norm.py](#)

[Liger-Kernel / src / liger_kernel / ops / jsd.py](#)

[Liger-Kernel / src / liger_kernel / ops / swiglu.py](#)

[Liger-Kernel / src / liger_kernel / ops / geglu.py](#)

[triton / python / tutorials / 01-vector-add.py](#)

[triton / python / tutorials / 03-matrix-multiplication.py](#)

torch.addmm

Performance

Can Decoupled Triton generate Triton kernels with performance that is **at least** equivalent to that of manually written Triton kernels?

Yes, with the existing scheduling directives Decoupled Triton can generate competitive Triton kernels for many common ML ops.

[Liger-Kernel / src / liger_kernel / ops / rms_norm.py](#)

[triton / python / tutorials / 01-vector-add.py](#)

[Liger-Kernel / src / liger_kernel / ops / layer_norm.py](#)

[triton / python / tutorials / 03-matrix-multiplication.py](#)

[Liger-Kernel / src / liger_kernel / ops / jsd.py](#)

torch.addmm

[Liger-Kernel / src / liger_kernel / ops / swiglu.py](#)

torch.mm

[Liger-Kernel / src / liger_kernel / ops / geglu.py](#)

Performance

Can Decoupled Triton generate Triton kernels with performance that is **at least** equivalent to that of manually written Triton kernels?

Yes, with the existing scheduling directives Decoupled Triton can generate competitive Triton kernels for many common ML ops.

More research directions from here ...

Autotuner

Autotuner

Decoupled Triton

```
Func f;  
Var x;  
Var y;  
  
f[x, y] = ...
```

```
f.compile_to_kernel();
```

Autotuner

Decoupled Triton

```
Func f;  
Var x;  
Var y;  
  
f[x, y] = ...
```

```
f.autotune(...);  
f.compile_to_kernel();
```

Autotuner

Decoupled Triton

```
Func f;  
Var x;  
Var y;  
  
f[x, y] = ...
```

```
f.autotune(...);  
f.compile_to_kernel();
```

Autotuner to search the space of schedules automatically to determine the best schedule for a given hardware.

Backward Pass Kernels

Backward Pass Kernels

Decoupled Triton

```
Func f;  
Var x;  
Var y;  
  
f[x, y] = ...  
  
f.compile_fwd();
```

Triton Python

```
import triton  
import triton.language as tl  
  
@triton.jit  
def f_fwd_kernel(...):  
    ...
```

Backward Pass Kernels

Decoupled Triton

```
Func f;  
Var x;  
Var y;  
  
f[x, y] = ...  
  
f.compile_fwd();  
f.compile_bwd();
```

Triton Python

```
import triton  
import triton.language as tl  
  
@triton.jit  
def f_fwd_kernel(...):  
    ...  
  
@triton.jit  
def f_bwd_kernel(...):  
    ...
```

Backward Pass Kernels

Decoupled Triton

```
Func f;  
Var x;  
Var y;  
  
f[x, y] = ...  
  
f.compile_fwd();  
f.compile_bwd();
```

Automatically generate the corresponding Triton kernel for the backward pass.

Triton Python

```
import triton  
import triton.language as tl  
  
@triton.jit  
def f_fwd_kernel(...):  
    ...  
  
@triton.jit  
def f_bwd_kernel(...):  
    ...
```

Backward Pass Kernels

Decoupled Triton

```
Func f;  
Var x;  
Var y;  
  
f[x, y] = ...  
  
f.compile_fwd();  
f.compile_bwd();
```

Automatically generate the corresponding Triton kernel for the backward pass.

It will need different scheduling.

Triton Python

```
import triton  
import triton.language as tl  
  
@triton.jit  
def f_fwd_kernel(...):  
    ...  
  
@triton.jit  
def f_bwd_kernel(...):  
    ...
```

User Defined Program ID Remapping

User Defined Program ID Remapping

`group()`, `dilate()`, and `aggregate_and_sequentialize()` are patterns for remapping program ID to result block.

User Defined Program ID Remapping

`group()`, `dilate()`, and `aggregate_and_sequentialize()` are patterns for remapping program ID to result block.

Can we create a system to allow the user to define a mapping function that maps program ID to result block?

User Defined Program ID Remapping

`group()`, `dilate()`, and `aggregate_and_sequentialize()` are patterns for remapping program ID to result block.

Can we create a system to allow the user to define a mapping function that maps program ID to result block?

Would such a system be useful?

Program Synchronization

Program Synchronization

Program synchronization to allow reductions of values across different programs.